

Blockchain Applications

Lecture 4: Stablecoins I – Measuring Depeg Risk and On-Chain Early Warning



Daniel Traian PELE

Bucharest University of Economic Studies
IDA Institute Digital Assets
Blockchain Research Center
Master Blockchain for Business (MBB), FABIZ
24 November 2026

Opening: How Far and How Often Does a Stablecoin Miss?



Richard K. Lyons. **Photo: Haas School of Business, UC Berkeley (2008); CC BY-SA 3.0; Wikimedia Commons**

"Our central thesis is that decentralization of issuance and access to arbitrage trades with the issuing treasury are key factors determining the efficiency of the peg."
— *Lyons & Viswanath-Natraj (2023)*

- Daily deviations of USDT, USDC and DAI from 1 USD average within ± 0.5 basis points (bp; 1 bp = 0.01%)
  - the risk is in rare episodes: USDC/USDT, hourly since 2022, went beyond 15 bp in 30 episodes (an episode starts after 24 quiet hours)
- Silicon Valley Bank (SVB) weekend, 11 March 2023: USDC fell to $-1\,036$ bp
  - 211 bp below its redemption value had the 3.3 billion USD of reserves held at SVB been lost entirely (-825 bp)
- DAI was beyond 15 bp on 26.6% of days in 2021 and on 0.3% in 2026; USDT on 3.0% in 2026
- An early-warning alarm built on pool data detects 72% of test events 14 hours ahead, with 2.5 false alerts a month

Central chart: The Three Dollars, Day by Day

Today's Question and Route

- **Question**: a stablecoin promises one dollar. How far and how often does it miss, and can we see a miss coming?
  - stablecoins settle most crypto trades and serve as collateral in decentralised finance (DeFi): a deviation of a few basis points can trigger liquidations
- **Route**
  - measuring the peg: deviations in basis points, tails, episodes; the Silicon Valley Bank (SVB) weekend of March 2023
  - where the price is formed: automated market maker (AMM) invariants: constant product and StableSwap
  - early warning: event-level evaluation, utility, results of a recent working paper (Chaffard, Găman, Staures, Ginavar, Jheng, Härdle & Pele (2026))
- Data: EODHD daily (2021–2026) and hourly (2022–2026) prices of USDT, USDC, DAI; DefiLlama supply (21 October 2026)

From Lecture 3: One Depeg, Three Measurements

- Lecture 3 summarised risk by one daily quantile: VaR 1% $= -q_{0.01}$, a positive loss
  - USDC, daily closes 2021–2026: $q_{0.01} = -7.6$ basis points (bp), so VaR 1% $= 7.6$ bp
- The SVB weekend (March 2023), measured three ways on hourly USDC prices
  - magnitude: lowest hourly close $-1\,036$ bp (11 March, 07:00 UTC); lowest daily close -285 bp
  - duration: 62 hours from the first hour at or below -15 bp to the return above -15 bp; 44 hours at or below -100 bp
  - event probability: how often such episodes start, and whether the next 24 hours will contain one
- A daily VaR answers none of the three
  - it is one quantile of daily closes: silent on the depth beyond it, on how long a depeg lasts and on when the next one starts
  - this lecture measures all three: deviations and tails, episodes, and an early-warning alarm

Learning Outcomes

- Measure a peg deviation in basis points and describe its distribution, tails and episodes
- Explain price impact under constant-product and StableSwap AMMs
- Explain why pool composition reveals liquidity stress before the price
- Evaluate an early-warning system with event-level recall, lead time, false alerts and utility

Abbreviations (1/2)

AMM automated market maker
AR autoregressive (model); AR(1): first order
ARIMA autoregressive integrated moving average
AUC area under the receiver operating characteristic curve
AUPRC area under the precision–recall curve
bn billion
bp basis point (0.01%)
BTC Bitcoin
CDF cumulative distribution function
CI confidence interval
CRPS continuous ranked probability score
DAI Maker/Sky stablecoin (crypto-backed)
DeFi decentralised finance
EODHD EOD Historical Data (data vendor)
ES expected shortfall
ETH Ether (Ethereum)
FABIZ Faculty of Business Administration in Foreign Languages, ASE
FDIC Federal Deposit Insurance Corporation
Fed Federal Reserve
FPCA functional principal component analysis
FRM Financial Risk Meter
GARCH generalised autoregressive conditional heteroskedasticity
GARCH-t GARCH with Student-t errors
GPD generalised Pareto distribution
HC heteroskedasticity-consistent
HHI Herfindahl–Hirschman index

Abbreviations (2/2)

JIT just-in-time (liquidity)
KS Kolmogorov–Smirnov (test)
KYC know your customer
m million
OLS ordinary least squares
PIT probability integral transform
PYUSD PayPal USD
QL quantile (pinball) loss
sd standard deviation
SE standard error
SHAP SHapley Additive exPlanations
SSR sum of squared residuals
SVB Silicon Valley Bank
TAR threshold autoregression
TiDE Time-series Dense Encoder
twCRPS threshold-weighted CRPS
USD US dollar
USD1 World Liberty Financial USD
USDC USD Coin (Circle)
USDe Ethena USDe
USDG Global Dollar (Paxos)
USDS Sky Dollar
USDT Tether
UTC Coordinated Universal Time
V3 version 3 (Uniswap V3)
VaR value-at-risk

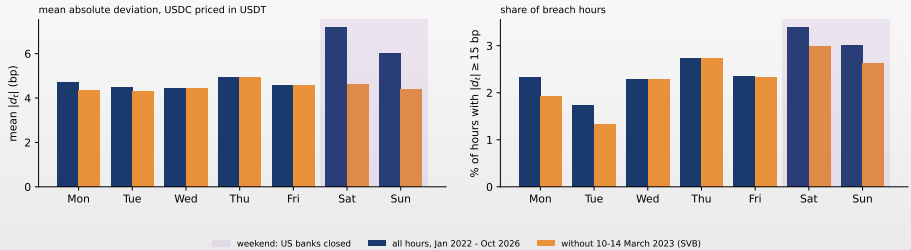
Reading and Tools

- Main paper: Chaffard, Găman, Staures, Ginavar, Jheng, Härdle & Pele (2026), *Onchain Monitoring of Stablecoin Depeg Risk*, working paper
- Data of the paper: Chaffard (2026), *Stablecoin DeFi Hourly Dataset*, Zenodo, the hourly on-chain data set behind its results
- AMMs: Egorov (2019) (StableSwap), Adams et al. (2021) (Uniswap V3)
- Threshold models: Hansen (1996); forecast evaluation: Gneiting & Raftery (2007), Gneiting & Ranjan (2011); attribution: Lundberg & Lee (2017)
- Market data: EODHD stablecoin prices, DefiLlama stablecoin supply; tools: Python (`numpy`, `scipy`, `scikit-learn`)
- Next lecture (Stablecoins II): the De-pegging Pressure Index and the Financial Risk Meter (FRM), across many stablecoins

Mechanisms

- A stablecoin is a token that targets a fixed price P^* in a reference asset, here $P^* = 1$ USD
- **Fiat-backed** (USDT, USDC): reserves in cash and Treasury bills; the issuer mints and redeems at par for eligible clients
- **Crypto-backed** (DAI/USDS): over-collateralised loans in a smart contract; liquidations protect the peg
- **Algorithmic** (TerraUSD, May 2022): a mint/burn swap against a volatile token; no external collateral
- The peg is held by **arbitrage**: if $P_t < 1$, buy on the market and redeem at 1; if $P_t > 1$, mint at 1 and sell
  - arbitrage needs access (know-your-customer (KYC) rules, minimum size), time (bank hours) and trust in the reserves
  - when one of the three fails, the market price can leave the peg

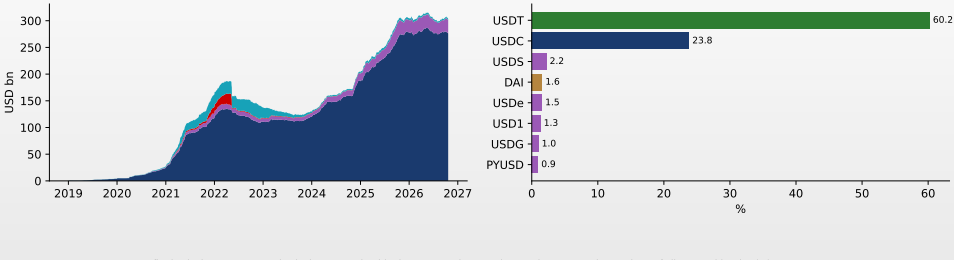
When Banks Are Closed



- USDC priced in USDT, hourly, January 2022 – 23 October 2026: mean $|d_t|$ 6.6 bp on weekend hours, 4.6 bp on weekday hours
- Without 10–14 March 2023 (SVB): 4.5 bp on both; one weekend makes the whole difference in the mean
- Breach hours ($|d_t| \geq 15$ bp), SVB days left out: 2.8% of weekend hours, 2.1% of weekday hours
- Redemption through the issuer pauses at weekends: the typical deviation is unchanged, breaches are a little more frequent; hours are autocorrelated, so this is a comparison of shares, not a test
- Source: own computation on EODHD hourly prices of USDC and USDT, January 2022 – 23 October 2026

BA_L4_weekday

Who Issues the Dollars on Chain



- Supply of USD stablecoins: 4.2 bn USD (Jan 2020), 187.1 bn (May 2022), 129.8 bn (Jan 2024), 303.4 bn (21 Oct 2026)
- Fiat-backed: 90.9%, crypto-backed 8.1%; algorithmic designs vanished after TerraUSD
- Two issuers hold 84% (USDT 60.2%, USDC 23.8%): a USDC/USDT dislocation is a system-level event

BA_L4_supply ▶ Interactive

Three Distinctions

- **Fundamental vs market-defined peg**
  - fundamental: the redemption value; market-defined: the price on a given venue
  - a token has no single price: on 10 October 2025 USDe traded at 0.65 USD on one exchange while other venues stayed near 1
- **Exogenous vs endogenous**
  - exogenous: reserve losses, regulation, sanctions; arrive as news, hard to predict
  - endogenous: liquidity and order flow inside the market; visible on chain *before* the price moves
- **Transient vs persistent**
  - transient: hours, closed by arbitrage; persistent: a loss of confidence (TerraUSD)
  - a transient depeg is not harmless: liquidations at a bad price are irreversible
- This lecture: market-defined, mostly endogenous, often transient risk, where monitoring is feasible

Definitions

- Deviation in basis points (1 bp = 0.01%):

$$d_t = 10^4 \, (P_t - 1), \qquad d_t^{\mathrm{USDC/USDT}} = 10^4 \left(\frac{P_t^{\mathrm{USDC}}}{P_t^{\mathrm{USDT}}} - 1 \right)$$

  - closing price vs intraday low: the low is the lowest price *reported* in the interval, aggregated over venues
  - whether a given amount could trade at that price depends on the venue, its depth and the trade size

- Exceedance probability and breach indicator, threshold δ (here 15 bp):

$$\hat{P}(|d| \geq x) = \frac{1}{T} \sum_{t=1}^{T} \mathbb{1}\{|d_t| \geq x\}, \qquad b_t = \mathbb{1}\{|d_t| \geq \delta\}$$

- An **episode** starts at a breach hour preceded by 24 *observed* hours without a breach
  - hours inside an episode are not independent observations: count episodes, not hours
  - a missing hour is never counted as quiet: a breach after a gap is ambiguous and is set aside

Worked Example: Finding an Episode Start

- Hourly EODHD closes on 29 June 2026, UTC: USDC 0.999554 and USDT 0.998109 at 14:00; USDC 0.999712 and USDT 0.998199 at 15:00

$$d_{14:00} = 10^4 \left(\frac{0.999554}{0.998109} - 1 \right) = 14.5 \text{ bp}, \qquad d_{15:00} = 10^4 \left(\frac{0.999712}{0.998199} - 1 \right) = 15.2 \text{ bp}$$

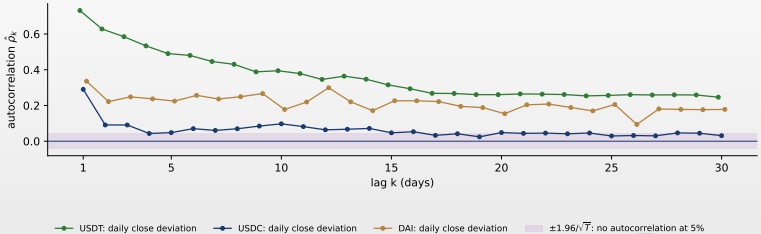
- Breach indicator with $\delta = 15$ bp: $b_{14:00} = 0$, $b_{15:00} = 1$
- Quiet spell: the 24 hours from 28 June, 15:00 to 29 June, 14:00 are all observed; the largest $|d_t|$ among them is 14.7 bp
  - so 15:00 starts an episode; at 16:00 $d_t = 13.8$ bp, and the episode lasts one hour
- Contrast: 4 February 2026, 17:00, $d_t = 15.9$ bp; the previous 24 hours show no breach, but one of them is missing
  - an ambiguous onset: the missing hour could have been a breach, so this hour is not counted as an episode

Peg Deviations, Daily, 2021–2026

bp	mean	sd	$q_{0.01}$	$q_{0.99}$	min (date)	% $\lvert d\rvert \geq 15$	% $\lvert d\rvert \geq 50$	min low	ex. kurt	$\hat{\rho}_1$
USDT	0.5	7.3	−15.5	21.3	−41 (11 May 22)	3.6	0.14	−515	40	0.73
USDC	−0.2	7.4	−7.6	9.7	−285 (11 Mar 23)	0.7	0.14	−1226	1045	0.29
DAI	0.4	10.8	−17.2	36.6	−261 (11 Mar 23)	5.2	0.57	−1030	185	0.34

- 2 122 daily closes, EODHD aggregated prices; earlier years are excluded because the aggregate includes illiquid venues
- Means lie within ± 0.5 bp: the peg holds *on average*; the risk is in the tails
- USDC: excess kurtosis 1 045 (0 for the Normal distribution); one weekend dominates; its 1% quantile (-7.6 bp) is the closest to zero of the three
- USDT: smallest extreme, but most persistent (lag-1 autocorrelation $\hat{\rho}_1 = 0.73$) and most often beyond 15 bp after 2022
- The lowest reported intraday prices are about 4 to 12 times deeper than the worst closes; whether a liquidation used them depends on the protocol's price oracle
- Source: own computation on EODHD daily closes and intraday lows of USDT, USDC and DAI, 2021 – 23 October 2026

How Long Does a Deviation Last?



- Autocorrelation $\hat{\rho}_k$ of the daily close deviation at lags of 1 to 30 days; the band $\pm 1.96/\sqrt{T} = \pm 0.043$ ($T = 2\,122$ days) marks no autocorrelation at the 5% level
- USDT: $\hat{\rho}_{\mathbf{1}} = 0.73$, $\hat{\rho}_{\mathbf{7}} = 0.45$, $\hat{\rho}_{\mathbf{30}} = 0.25$: a discount lasts for weeks
- USDC: $\hat{\rho}_{\mathbf{1}} = 0.29$, $\hat{\rho}_{\mathbf{2}} = 0.09$: a deviation closes within a day or two
- DAI: $\hat{\rho}_{\mathbf{1}} = 0.34$, then a plateau near 0.2 up to 30 days
- An AR(1) with $\rho = 0.73$ has a half-life of $\ln 0.5 / \ln 0.73 = 2.2$ days and $\rho^{\mathbf{30}} \approx 0.0001$; the observed 0.25 at 30 days points to slow shifts in the level of USDT
- Source: own computation on EODHD daily closes of USDT, USDC and DAI, 2021 – 23 October 2026

BA_L4_acf

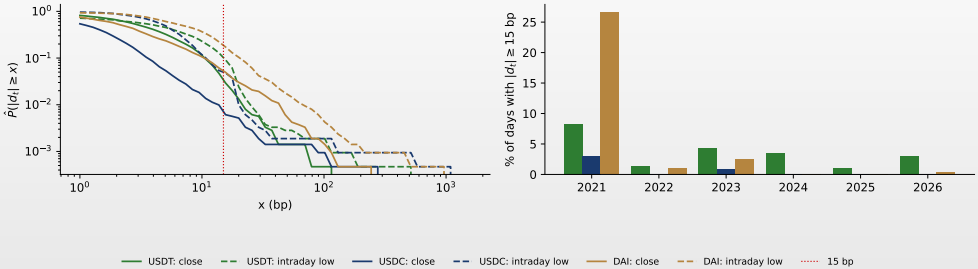
The Three Dollars, Day by Day



- Closes stay inside ±15 bp most of the time; the lows (red) reveal the stress days: May 2022 (Terra), March 2023 (SVB)

BA_L4_dev_daily ▶ Interactive

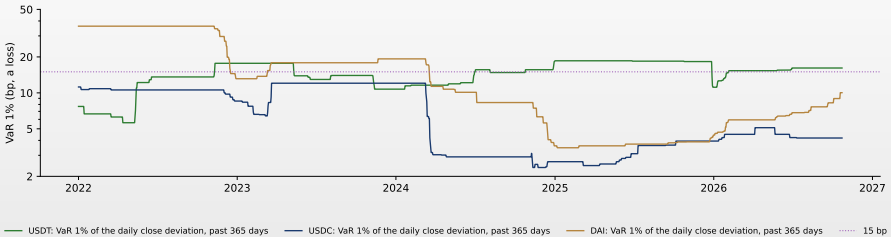
Tails and Their Frequency



- Left: on a log-log scale the exceedance curves decay slowly beyond 15 bp: heavy tails (Lecture 2); a power law would need an explicit tail-index estimate and test
- Right: DAI was beyond 15 bp on 26.6% of days in 2021, 0.3% in 2026; USDT is the only one still breaching regularly (3.0% in 2026)

BA_L4_exceed ▶ Interactive

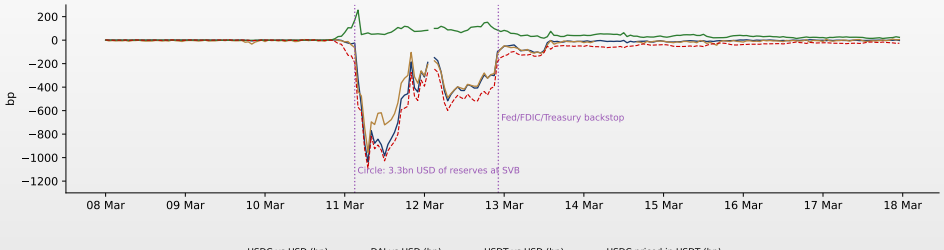
Tail Risk Moves: Rolling VaR 1%



- VaR 1% $= -q_{0.01}$ of the daily close deviation over the previous 365 days, a positive loss in bp (Lecture 3); full sample 2021–2026: USDT 15.5, USDC 7.6, DAI 17.2 bp
- USDT: from 5.6 bp (April 2022) to 18.6 bp (1 January 2025); 16.2 bp on 23 October 2026: the only coin whose tail widened
- USDC: 12.0 bp for a year after SVB, 4.2 bp now; DAI: 36.2 bp at the start of 2022, 3.5 bp in January 2025, 10.0 bp now
- With 365 days, the 1% quantile lies between the 4th and 5th lowest close: the line moves in steps when single days enter or leave the window
- Source: own computation on EODHD daily closes of USDT, USDC and DAI, 2021 – 23 October 2026

BA_L4_rollvar ▶ Interactive

The SVB Weekend, Hour by Hour



- Hourly minimum: USDC −1 036 bp and DAI −955 bp, both at 07:00 Coordinated Universal Time (UTC) on Saturday 11 March 2023; USDT rose above 1 (flight to the other dollar)
- USDC back above −100 bp at 22:00 UTC on 12 March, just before the backstop; above −15 bp only on 13 March, 14:00 UTC
- DAI followed USDC because a large part of its collateral was USDC: contagion through the balance sheet
- Source: own computation on EODHD hourly prices of USDC, USDT and DAI, March 2023

BA_L4_svb ▶ Interactive

Worked Example: Was the SVB Discount Fair?

- Circle disclosed 3.3 bn USD of about 40 bn USD reserves at SVB: share $s = 3.3/40 = 8.25\%$
- With recovery rate R on the SVB deposit, the redemption value per token is

$$V(R) = 1 - s\,(1 - R), \qquad V(0) = 0.9175 \;\Rightarrow\; d = -825 \text{ bp}$$

- The market reached $-1\,036$ bp, 211 bp below the value under *zero* recovery
- The 211 bp are the discount *unexplained by this reserve-loss scenario*, not an estimate of a liquidity premium
- Possible components, not separated by this calculation
  - time: banks were closed over the weekend, so redemption at par was impossible
  - uncertainty: about the size of the loss and about the other reserve holdings
  - frictions: arbitrage capital, client access to redemption, fragmented venues
- A peg can fail with reserves that cover 92% or more: access matters as well as backing

Constant-Product AMM

- A pool holds x of token X and y of token Y; trades keep the invariant $x\,y = k$
- Marginal price of X in Y: $p = y/x$. Selling Δx returns

$$\Delta y = y - \frac{k}{x + \Delta x} = \frac{y\,\Delta x}{x + \Delta x}, \qquad \text{slippage} = 1 - \frac{\Delta y}{\Delta x} = \frac{\Delta x}{x + \Delta x} \quad (p = 1)$$

- Worked example: pool 10 m X + 10 m Y, sell 1 m X

$$\Delta y = \frac{10 \cdot 1}{11} = 0.909 \text{ m}, \qquad \text{slippage } 9.1\%$$

- Liquidity is spread over all prices from 0 to ∞: robust but very expensive for stable pairs

StableSwap (Curve): Flat Near Parity

- Invariant for n coins, amplification A (Egorov (2019)):

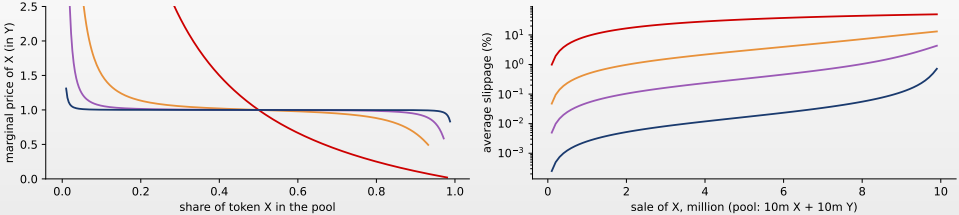
$$A\,n^n \sum_i x_i + D = A\,D\,n^n + \frac{D^{n+1}}{n^n \prod_i x_i}$$

  - x_i: balance of coin i; D: total balance when the pool is balanced at price 1 ($x_i = D/n$)
  - $A \to 0$: constant product; $A \to \infty$: constant sum ($x + y = D$, price exactly 1)

Sale of X (pool 10 m + 10 m)	constant product	$A = 10$	$A = 100$	$A = 2000$
0.1 m	0.99%	0.048%	0.005%	0.000%
1 m	9.09%	0.48%	0.05%	0.003%
5 m	33.3%	2.95%	0.33%	0.017%
8 m	44.4%	7.22%	1.04%	0.055%
9.5 m	48.7%	11.6%	2.94%	0.232%

- Average slippage: StableSwap is thousands of times cheaper near balance, but slippage rises steeply when the pool empties of one coin
- Source: own computation from the constant-product and StableSwap invariants; no market data

Prices Are Inelastic until They Are Not



- With $A = 100$ the price of X is 0.986 when X is 80% of the pool, 0.943 at 90%, 0.807 at 95%; with $A = 2000$: 0.999, 0.997, 0.988
- The pool absorbs one-sided flow with almost no price change, so the *price* hides the stress; the *composition* shows it

BA_L4_amm ▶ Interactive

From Pool State to a Warning: the Set-up

- Each hour t, features x_t are computed from data up to t only; three examples
  - Curve 3pool balance: entropy of the DAI/USDC/USDT shares ($\ln 3 = 1.0986$ when balanced)
  - Uniswap V3 depth: liquidity within a few ticks of the current price
  - recent deviation: $|d_t|$ and its maximum over the last 24 hours
- Event: a breach $|d_{t+h}| \geq 15$ bp within the next 24 hours gives the label $y_t \in \{0, 1\}$
- Time order: fit on the early years, choose the alert threshold on a later validation year, report once on the final test period
- Metrics, before the results
  - precision: share of alerted hours that are followed by a breach; recall: share of positive hours that are alerted
  - AUPRC: area under the precision–recall curve over all thresholds; a random score gets the base rate (5.9% in the paper), a perfect one 1
- The other features (Gegenbauer scores, ownership, Aave) summarise the same pools; definitions in the appendix

Data and Target in the Paper

- Chaffard, Găman, Staures, Ginavar, Jheng, Härdle & Pele (2026): Uniswap V3 USDC/USDT pool (0.01% fee), hourly, January 2022 – May 2026, 38 045 observations
- Feature blocks: pool price and flows; Curve 3pool shares and entropy; swap-impact curves; Gegenbauer scores; ownership (JIT share, Herfindahl–Hirschman index (HHI), width, age of positions); Aave utilisation, rates, liquidations; BTC/ETH volatility; a USD index from Chainlink; fear and greed index
- Target: a breach of $\delta = 15$ bp within the next 24 hours

$$y_t = \mathbb{1}\Big\{ \max_{1 \le h \le 24} |d_{t+h}| \ge \delta \Big\}$$

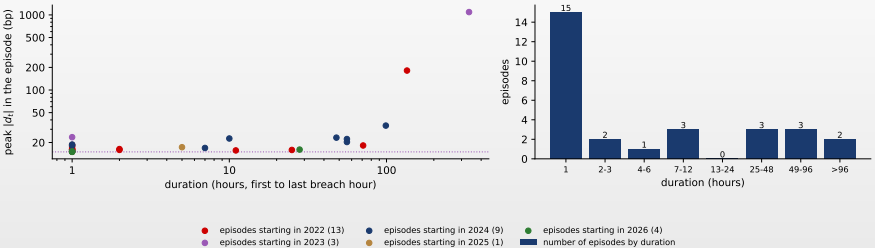
- Labels are not events: every hour in the 24 hours before a breach, and inside it, is labelled 1
  - Uniswap V3 pool in the paper: 2 227 positive hours (5.9%) but only 107 events
  - EODHD cross rate USDC/USDT, Jan 2022 – Oct 2026: 2 291 positive hours (7.1%), 30 episodes
- The event count depends on where the price is read: one pool breaches more often than a price aggregated over venues (107 vs 30 events; the periods also differ)

Episodes by Threshold

δ (bp)	episodes	ambiguous	breach hours	2022	2023	2024	2025	2026
10	98	21	4 282	30	7	31	11	19
15	30	20	1 044	13	3	9	1	4
25	5	3	358	3	1	1	0	0
50	2	1	105	1	1	0	0	0

- Episode: breach hour after 24 observed quiet hours; ambiguous: no breach observed in the previous 24 hours, but some of them missing (not counted)
- From 10 to 15 bp the number of episodes falls by about 70% (98 to 30); at 25 bp or more there are only 5
- The threshold is a policy choice: low δ gives many, mostly harmless events; high δ leaves too few events to learn or to test
- Source: own computation on EODHD hourly prices of USDC and USDT, January 2022 – 23 October 2026; 2.8% of hours missing

Episodes: Mostly Short and Shallow



- The 30 episodes at 15 bp; duration: hours from the first to the last breach hour before 24 hours without a breach
- 15 of 30 episodes last one hour, 18 at most 6 hours; median duration 1.5 hours; median peak $|d_t|$ 16.6 bp; 27 of 30 peak below 25 bp
- Two exceptions: the collapse of the FTX exchange, November 2022 (135 hours, peak 182 bp), and SVB, March 2023 (335 hours, peak 1 091 bp)
- Longer episodes go deeper: Spearman rank correlation of duration and peak 0.57
- Source: own computation on EODHD hourly prices of USDC and USDT, January 2022 – 23 October 2026

BA_L4_episodes ▶ Interactive

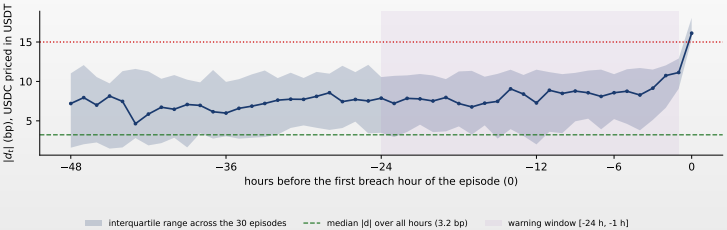
Worked Example: From Frequencies to Waiting Times

- Daily closes, 2021–2026: USDT beyond 15 bp on 3.6% of days, so $0.036 \times 365 = 13$ days a year
  - beyond 50 bp on 0.14% of days (3 of 2 122): once every $1/(0.00141 \times 365) = 1.9$ years
- Hourly USDC/USDT: 30 episodes in 4.81 years, a rate $\lambda = 30/4.81 = 6.2$ a year, one every $365/6.2 = 59$ days
- If episodes arrived independently at this rate (a Poisson process),

$$P(\text{at least one episode in a month}) = 1 - e^{-6.2/12} = 1 - e^{-0.52} = 0.41$$

- The data disagree with a constant rate
  - median gap between episodes 25 days; 7 of the 29 gaps are shorter than a week
  - only 5 of the 30 episodes fall in 2025–2026: episodes cluster in stress years
- A constant rate overstates risk in calm years and understates it in stress years

What the Price Shows before an Episode



- $|d_t|$ of USDC priced in USDT in the 48 hours before the first breach hour of each of the 30 episodes; median over all hours: 3.2 bp
- Median 24 hours before the breach: 7.9 bp, more than twice the typical level; 1 hour before: 11.1 bp
- In 25 of 30 episodes $|d_t|$ reached 10 bp or more somewhere in the warning window $[-24\,\text{h}, -1\,\text{h}]$
- A price rule can see many episodes coming, but 10–14 bp is also common without any breach (October 2026: 206 hours above 10 bp, none above 15): hence the false alerts below
- Source: own computation on EODHD hourly prices of USDC and USDT, January 2022 – 23 October 2026

BA_L4_preevent

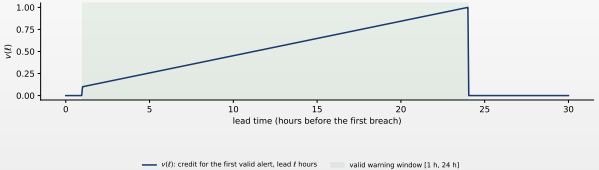
Evaluating an Alarm, Not a Classifier

- Row-level area under the receiver operating characteristic curve (AUC) and under the precision–recall curve (AUPRC) reward correct labels *during* an episode; an alert after the breach is no warning
- Event-level protocol
  - alerts: hours with score above a threshold; a 24-hour refractory period merges them into alert episodes
  - event e at time τ_e is detected if an alert falls in $[\tau_e - 24\text{h}, \tau_e - 1\text{h}]$; lead ℓ_e = hours before τ_e
  - all other alert episodes are false; F_m = false alerts per month
- Utility

$$U = \frac{1}{|E|} \sum_{e \in E} v(\ell_e) - 0.05\, F_m, \qquad v(\ell) = 0.10 + 0.90\, \frac{\ell - 1}{23} \ \text{ for } 1 \le \ell \le 24, \ \text{else } 0$$

- Alert threshold chosen on the preceding validation fold, maximising U subject to $F_m \le 2$ per month *on validation*; never re-tuned on the test fold, where F_m can exceed 2

Worked Example: Computing the Utility



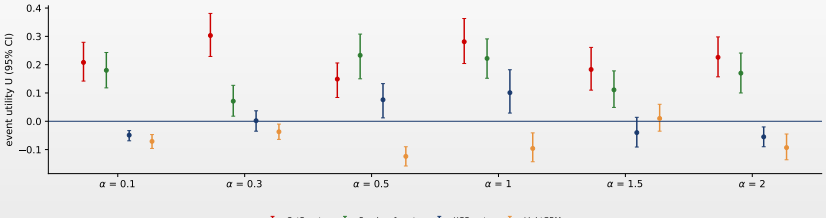
- Four events; first valid alerts 14 h, 2 h, none, 20 h before the breach; 2.5 false alerts per month

$$v = (0.609,\ 0.139,\ 0,\ 0.843), \qquad \bar{v} = 0.398, \qquad U = 0.398 - 0.05 \cdot 2.5 = 0.273$$

- One more false alert per month costs 0.05, about the value of detecting one extra event (out of four) with a 2-hour lead

BA_L4_utility ▶ Interactive

Results: Four Classifier Families



	U [95% CI]	recall	F_m	median lead (h)	AUC	AUPRC
CatBoost, $\alpha = 0.3$	0.303 [0.229, 0.381]	0.723	2.50	14.4	0.862	0.241
Random forest, $\alpha = 0.1$	0.180 [0.118, 0.243]	0.410	0.97	12.0	0.901	0.239

- Test fold: CatBoost detects 72% of events, 14 hours ahead, with 2.5 false alerts a month (the budget of 2 was imposed on validation); the random forest with the *best AUC* detects only 41%: AUC is not the right criterion
- Source: Chaffard et al. (2026), working paper, Table 1

BA_L4_paper_t1 ▶ Interactive

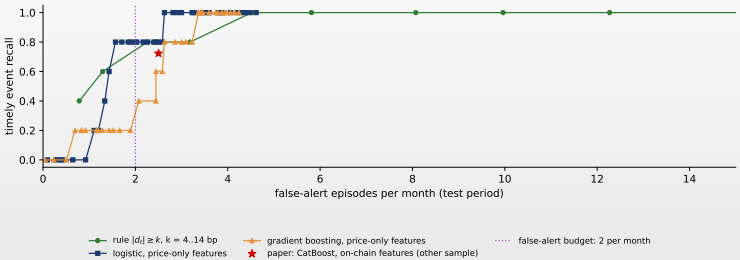
A Price-Only Baseline

- Question: can the price alone warn of a breach? The same alarm protocol, on the EODHD cross rate USDC/USDT
- Time order: train 2022–2023, choose the threshold in 2024 ($F_m \leq 2$), test 2 Jan 2025 – 23 Oct 2026; 24-hour gaps (embargoes) so that no label spans two periods
- Features: d_t, $|d_t|$, 24 h max and volatility of d, deviation from the weekly mean, DAI and USDT vs USD, BTC 24 h return and volatility

Test, 5 events	detected	recall	lead (h)	F_m	U	AUC	AUPRC
Logistic regression	0	0	–	0.09	−0.005	0.929	0.238
Gradient boosting	0	0	–	0.05	−0.002	0.880	0.267
Rule $\|d_t\| \geq k$, $k = 20$ chosen in 2024	0	0	–	0.18	−0.009	0.915	0.345

- AUC 0.88–0.93, yet no timely warning: high AUC comes from labelling hours *inside* episodes (base rate 1.5%)
- In 2024 every rule that detected events produced at least 2.6 false alerts a month, so the budget forced the threshold above 15 bp
- Comparison with the paper is indirect: venue, period and number of events differ
- Source: own computation on EODHD hourly prices of USDC, USDT, DAI and BTC, January 2022 – 23 October 2026

The Recall–False-Alert Trade-off



- Choosing k on the test data is look-ahead: the rule $|d_t| \geq 11$ bp would have reached recall 0.8, median lead 20.5 h, $U = 0.42$, but with 3.2 false alerts a month
- 4 of 5 events: Clopper–Pearson 95% interval for recall $[0.28, 0.99]$. With so few events, the ranking of alarms is fragile
- Source: own computation on EODHD hourly prices, test period 2 January 2025 – 23 October 2026

BA_L4_ew_tradeoff ▶ Interactive

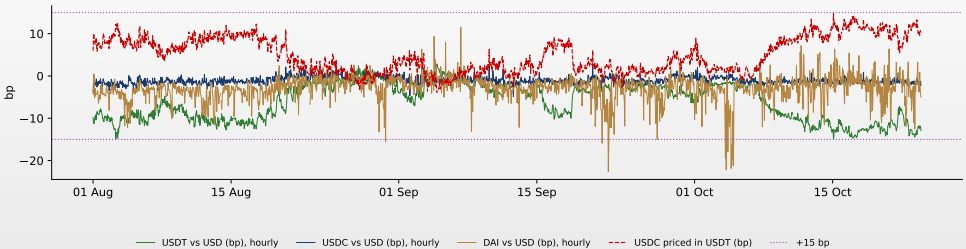
Worked Example: What Can Five Events Tell?

- Rule $|d_t| \geq 11$ bp on the test period: 4 of 5 events detected, recall 0.8
- Suppose the true recall is p; the number detected is $X \sim \text{Binomial}(5, p)$:

$$P(X \geq 4) = 5\,p^4(1-p) + p^5$$

  - $p = 0.5$: $5/32 + 1/32 = 6/32 = 0.19$
  - $p = 0.3$: $0.028 + 0.002 = 0.03$
  - $p = 0.723$ (the paper's CatBoost): 0.58
- An alarm with recall 0.5 scores 4 of 5 in about one test period out of five: five events cannot separate it from the paper's alarm
- The Clopper–Pearson 95% interval for 4 of 5, $[0.28, 0.99]$, says the same in one line
- To compare alarms, pool events across periods or venues, or report intervals next to every recall

The Last Twelve Weeks



- USDT below USD: −6.6 bp on average in August, −3.2 in September, −8.7 in October 2026
- USDC priced in USDT: October mean 7.5 bp, maximum 14.7 bp (15 October, 02:00 UTC); 206 of 552 hours above 10 bp, none above 15 bp
- A rule at 10 bp would have fired repeatedly; the paper's alarm would look at the pool composition and the liquidity curve before calling it stress
- Source: own computation on EODHD hourly prices of USDC and USDT, 1 August – 23 October 2026

BA_L4_recent ▶ Interactive

Summary

- Daily peg deviations average within ± 0.5 bp and have heavy tails; the risk is in rare episodes
  - USDC/USDT since 2022: 30 episodes beyond 15 bp, 2 beyond 50 bp
- SVB: USDC fell to $-1\,036$ bp, 211 bp below its zero-recovery value (-825 bp): reserve losses alone do not explain the discount
- AMMs keep prices flat until a pool is nearly one-sided: pool composition reveals stress before the price
- The paper's alarm detects 72% of test events 14 hours ahead, with 2.5 false alerts a month
- Price-only alarms on the aggregated USDC/USDT rate reach AUC up to 0.93 but meet the false-alert budget only by detecting none of the 5 test events

Interpretation Exercises

1. USDC has the 1% quantile closest to zero of the three coins (-7.6 bp) but the largest excess kurtosis (1 045). What do the two numbers say together about its risk?
2. A Curve pool with $A = 2000$ shows a price of 0.997 while 90% of it is USDC. Is USDC safe? What would you monitor instead of the price?
3. The random forest has AUC 0.901 and recall 0.41; CatBoost has AUC 0.862 and recall 0.72. Which would you deploy, and why does the ranking differ?
4. A rule $|d_t| \geq 11$ bp reaches $U = 0.42$ on the 2025–2026 test period. Give two reasons why this does not show that the rule is as good as the paper's model.

⊡ Next: Lecture 5, the De-pegging Pressure Index across stablecoins

Appendix: Further Material

Optional material for self-study.

- A Threshold Autoregression
- Estimates: Hourly USDC/USDT, 2022–2026
- The Fitted Band
- Worked Example: How Fast Does a Deviation Close?
- Uniswap V3: Concentrated Liquidity
- Curve 3pool: Entropy as an Imbalance Index
- A Fixed Basis for a Moving Curve
- Worked Example: Decomposing a Stylised Curve
- What Drives the Warnings? SHAP Attribution
- From an Alarm to a Distribution
- A Quantile Function with Pareto Tails
- What the Tail Parameters Do
- Scoring a Distribution: CRPS and twCRPS
- Worked Example: Two Forecasters, Three Outcomes
- Results: Distributional Accuracy
- Calibration and Ablation

A Threshold Autoregression

- If arbitrage costs c, small deviations are not worth trading; beyond c they are. A threshold autoregressive (AR) model captures this:

$$\Delta d_t = \mu + \rho_{in}\, d_{t-1}\, \mathbb{1}\{|d_{t-1}| \le \gamma\} + \rho_{out}\, d_{t-1}\, \mathbb{1}\{|d_{t-1}| > \gamma\} + \varepsilon_t$$

  - band hypothesis: $\rho_{in} \approx 0$, $\rho_{out} < 0$; half-life $h = \ln 0.5 / \ln(1+\rho)$
- Estimation: for each γ on a grid, ordinary least squares (OLS); $\hat{\gamma} = \arg\min_\gamma \mathrm{SSR}(\gamma)$, the sum of squared residuals (SSR)
- Test of a threshold effect, $H_0 : \rho_{in} = \rho_{out}$; γ is not identified under H_0, so the usual χ^2 fails; with a heteroskedasticity-consistent (HC) variance:

$$\sup W = \sup_{\gamma \in \Gamma} W(\gamma), \qquad W(\gamma) = \frac{(\hat{\rho}_{in} - \hat{\rho}_{out})^2}{\widehat{\mathrm{Var}}_{\mathrm{HC}}(\hat{\rho}_{in} - \hat{\rho}_{out})}$$

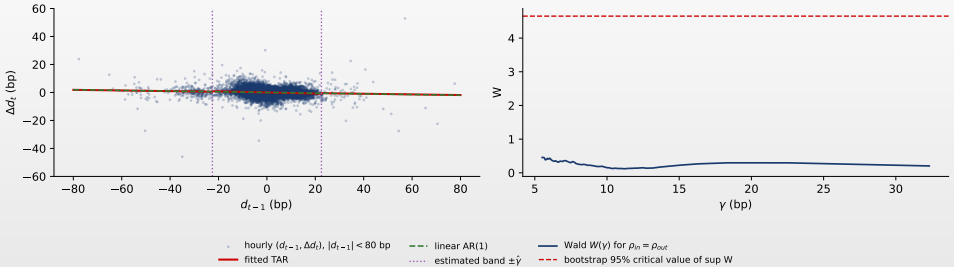
- p-value by the fixed-regressor bootstrap (Hansen (1996)): $y_t^* = \hat{e}_t z_t$, $z_t \sim N(0,1)$, recompute $\sup W$, $B = 199$

Estimates: Hourly USDC/USDT, 2022–2026

	coefficient	SE (HC)	half-life (hours)
$\hat{\gamma}$ (bp)	22.5		99.0% of hours inside
$\hat{\rho}_{in}$	−0.0396	0.0020	17.2
$\hat{\rho}_{out}$	−0.0233	0.0297	29.4
linear AR(1): $\hat{\rho}$	−0.0245		28.0
sup W (bootstrap 95% critical value)	0.5 (4.7),	p = 0.653	

- $n = 40\,545$ pairs of consecutive hours; SE: standard error
- No threshold effect: $\rho_{in} = \rho_{out}$ is not rejected; deviations decay in about a day either way
- The band story fails in this sample because the large deviations are SVB and Terra hours, when redemption was closed or trust was gone: arbitrage was not available, so reversion was *not* faster
- Lesson: a band model describes normal times; stress is a change of regime, not a larger deviation of the same process
- Source: own computation on EODHD hourly prices of USDC and USDT, January 2022 – 23 October 2026

The Fitted Band



- Left: the hourly change against last hour's deviation; slopes are small (≈ -0.04): reversion is slow at the hourly scale
- Right: the Wald profile $W(\gamma)$ never comes close to the bootstrap critical value

BA_L4_tar

Worked Example: How Fast Does a Deviation Close?

- Inside the band, $\hat{\rho}_{in} = -0.0396$: the expected deviation after h hours is

$$E[d_{t+h} \mid d_t] \approx (1 + \hat{\rho}_{in})^h \, d_t \quad (\hat{\mu} \approx 0)$$

- Start from $d_t = +12$ bp: USDC priced above USDT, i.e. USDT is the cheap coin, as on 23 October 2026

$$\text{24 hours: } 0.9604^{24} = 0.379, \qquad E[d_{t+24}] \approx 12 \times 0.379 = +4.55 \text{ bp}$$

- Half-life $\ln 0.5 / \ln 0.9604 = 17.2$ hours
- Interpretation
  - a deviation of this size is expected to shrink by about 60% in a day
  - if it does not, the process is not "inside the band": that persistence is itself a warning signal

Uniswap V3: Concentrated Liquidity

- Liquidity providers choose a price range; prices live on ticks $p_k = 1.0001^k$, one tick = 1 bp
- Active liquidity at tick j: $L_{t,j}$ = sum of all positions whose range contains j; the curve $j \mapsto L_{t,j}$ changes every hour
- Cumulative liquidity proxy from the current tick k_t (Chaffard, Găman, Staures, Ginavar, Jheng, Härdle & Pele (2026)), in liquidity units, not in tokens:

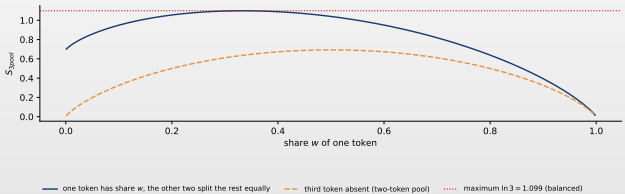
$$Q_t^+(k) = \sum_{j=k_t}^{k} L_{t,j} \;\; (k > k_t), \qquad Q_t^-(k) = \sum_{j=k}^{k_t} L_{t,j} \;\; (k < k_t)$$

  - token amounts to cross tick j (price $p = y/x$): up, pay $\Delta y_j = L_{t,j}(\sqrt{p_{j+1}} - \sqrt{p_j})$ of Y; down, pay $\Delta x_j = L_{t,j}(1/\sqrt{p_j} - 1/\sqrt{p_{j+1}})$ of X
  - area under the curves: global depth on each side; slope at k_t: local depth
  - a sudden flat part: a **liquidity cliff**, where a small trade moves the price a lot
- Data: The Graph subgraphs, hourly, ± 50 ticks around the peg; positions minted and burned in the same block are just-in-time (JIT) liquidity and may not be there under stress

Curve 3pool: Entropy as an Imbalance Index

- Normalised balances $w_{\mathrm{DAI}} + w_{\mathrm{USDC}} + w_{\mathrm{USDT}} = 1$:

$$S_{3pool} = -\sum_i w_i \ln w_i, \qquad 0 \le S \le \ln 3 = 1.0986$$



- Worked example: balanced $(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$: $S = 1.0986$; $(0.20, 0.40, 0.40)$: 1.0549; $(0.05, 0.10, 0.85)$: 0.5182
- Flat near balance, steep near a corner: small changes in S at high levels are noise, a fall below about 1 is stress

BA_L4_entropy

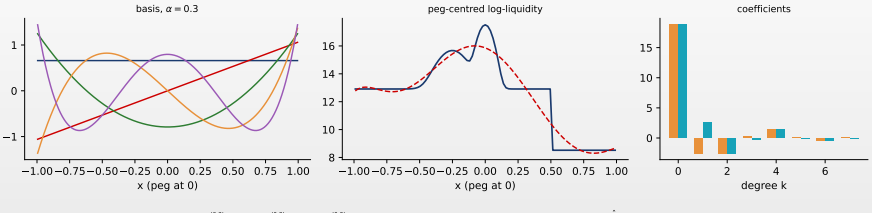
A Fixed Basis for a Moving Curve

- Peg-centred coordinate $x \in [-1, 1]$ (peg at $x = 0$), log-liquidity $\ell_t(x) = \log\{L_t(x) + 1\}$
- Expansion in Gegenbauer polynomials, orthonormal under the weight w_α:

$$\ell_t(x) \approx \sum_{k=0}^{K} \beta_{t,k}\, \phi_k^{(\alpha)}(x), \qquad w_\alpha(x) = (1 - x^2)^{\alpha - 1/2}, \qquad \beta_{t,k} = \int_{-1}^{1} \ell_t(x)\, \phi_k^{(\alpha)}(x)\, w_\alpha(x)\, dx$$

- Reading the coefficients
  - $k = 0$: level of liquidity; $k = 1$: left-right asymmetry; $k = 2$: concentration at the peg
  - higher k: local notches and cliffs; odd k asymmetric, even k symmetric
- Energies: $E_t^{odd} = \sum_{k \text{ odd}} \beta_{t,k}^2$, $E_t^{even} = \sum_{k \geq 2 \text{ even}} \beta_{t,k}^2$
- α moves the weight between centre and edges; it is tuned as a hyperparameter. A fixed basis keeps the scores comparable over time, unlike rolling functional principal component analysis (FPCA)

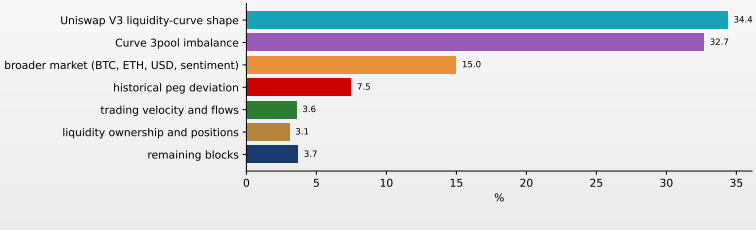
Worked Example: Decomposing a Stylised Curve



- Stylised curve: a deep band at the peg, extra depth below parity, a cliff above $x = 0.5$; $\alpha = 0.3$, $K = 7$
- $\hat{\beta} = (18.78, -2.52, -2.50, 0.32, 1.45, 0.08, -0.43, 0.07)$; $E^{odd} = 6.48$, $E^{even} = 8.55$
- The mirror image has the *same* even coefficients and odd coefficients with the opposite sign: the odd scores carry the direction of fragility
- Source: own computation on a stylised liquidity curve; no market data

BA_L4_gegenbauer

What Drives the Warnings? SHAP Attribution



- SHapley Additive exPlanations (SHAP; Lundberg & Lee (2017)): $f(\mathrm{x}) = \phi_{\mathbf{0}} + \sum_j \phi_j(\mathrm{x})$, shares of mean $|\phi_j|$ grouped by economic block (interventional, main effects)
- Past peg deviations explain only 7.5%: the model warns from liquidity shape (34.4%) and 3pool imbalance (32.7%), not from the price itself
- Attribution, not causality: correlated on-chain features share credit
- Source: Chaffard et al. (2026), working paper, SHAP shares of its CatBoost model on the test set

BA_L4_paper_shap

From an Alarm to a Distribution

- An alarm says *whether*; a risk manager also needs *how far*: the probability of crossing a liquidation level, the expected shortfall (ES) of a depeg
- The paper forecasts the full conditional distribution of d_{t+h}, $h = 1, \ldots, 24$
- Architecture, in brief
  - two branches: the target history (dilated causal convolutions) and the covariates (one token per variable)
  - reversible instance normalisation; a gated residual network scores covariates, sparsemax sets most weights to exactly zero
  - cross-attention from forecast horizons to the selected covariates; an output head returns a quantile function
- Benchmarks: naive, autoregressive integrated moving average (ARIMA), generalised autoregressive conditional heteroskedasticity with Student-t errors (GARCH-t), Time-series Dense Encoder (TiDE) and TimeXer; the two neural benchmarks get the same distributional head and loss
- Test period: January 2025 – May 2026

A Quantile Function with Pareto Tails

- Body: monotone I-spline, centred at the median and scaled

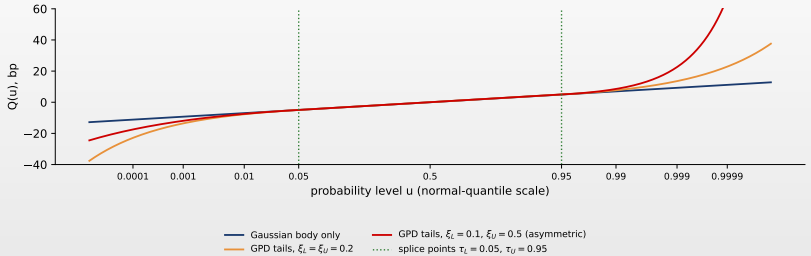
$$Q_t^0(u) = \sum_{k=1}^{K} a_{t,k} I_k(u),\ a_{t,k} > 0; \qquad Q_t^B(u) = \ell_t + \frac{s_t}{\sqrt{V_t}} \{Q_t^0(u) - Q_t^0(0.5)\}$$

- Tails: generalised Pareto distribution (GPD), joined smoothly at τ_L and τ_U

$$Q_t(u) = \begin{cases} z_L - \dfrac{\beta_L}{\xi_L}\left[\left(\frac{u}{\tau_L}\right)^{-\xi_L} - 1\right], & u < \tau_L \\ Q_t^B(u), & \tau_L \le u \le \tau_U \\ z_U + \dfrac{\beta_U}{\xi_U}\left[\left(\frac{1-u}{1-\tau_U}\right)^{-\xi_U} - 1\right], & u > \tau_U \end{cases}$$

- $z_L = Q_t^B(\tau_L)$, $\beta_L = \tau_L\, Q_t^{B\prime}(\tau_L)$ and symmetrically for the upper tail; separate ξ_L, ξ_U allow USDC and USDT depegs to differ
- For $\xi \ge 1/2$ the tail has infinite variance (Lecture 2: tail index $1/\xi$)

What the Tail Parameters Do



- Same Gaussian body (sd 3 bp), splice at 5% and 95%; on the normal-quantile scale the Gaussian is a straight line
- 99.9% quantile: Gaussian 9.3 bp; GPD with $\xi = 0.2$: 13.6 bp; with $\xi_U = 0.5$: 22.6 bp, while the lower tail ($\xi_L = 0.1$) stays at -11.9 bp
- Source: own illustration with assumed tail parameters; no market data

BA_L4_gpd

Scoring a Distribution: CRPS and twCRPS

- Continuous ranked probability score (Gneiting & Raftery (2007)), a proper scoring rule:

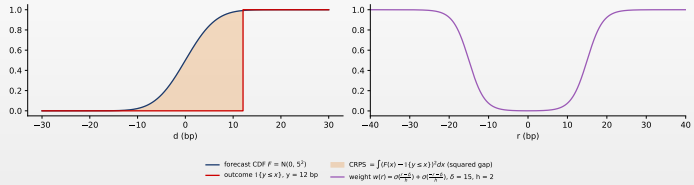
$$\mathrm{CRPS}(F, y) = \int_{-\infty}^{\infty} \big(F(x) - \mathbb{1}\{y \le x\}\big)^2 dx = 2\int_0^1 \rho_u\big(y - Q(u)\big)\, du$$

  - $\rho_u(e) = e\,(u - \mathbb{1}\{e < 0\})$: the pinball loss; CRPS averages it over all quantiles
  - Normal: $\mathrm{CRPS}(N(\mu, \sigma^2), y) = \sigma\big[z(2\Phi(z) - 1) + 2\varphi(z) - 1/\sqrt{\pi}\big]$, $z = (y - \mu)/\sigma$
- Threshold-weighted CRPS (Gneiting & Ranjan (2011)): put the weight where depegs are

$$\mathrm{twCRPS} = 2\int_0^1 \rho_u\big(v(y) - v(Q(u))\big)\, du, \quad v(r) = \int_0^r w(s)\, ds, \quad w(r) = \sigma\Big(\frac{r - \delta}{h}\Big) + \sigma\Big(\frac{-r - \delta}{h}\Big)$$

  - $\sigma(z) = 1/(1 + e^{-z})$, the logistic function; h sets how sharply the weight rises around $\pm\delta$
- The paper uses twCRPS as training loss and both scores for evaluation

Worked Example: Two Forecasters, Three Outcomes

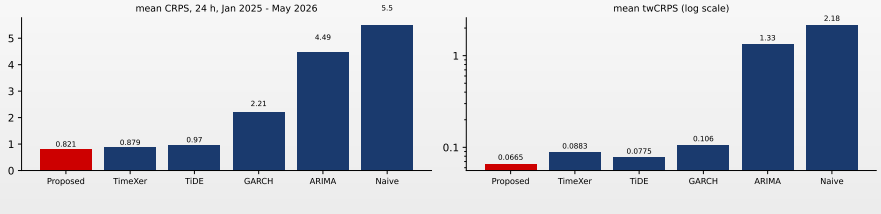


forecast	CRPS $y=1$	CRPS $y=12$	CRPS $y=25$	twCRPS ($\delta=15$, $h=2$, illustrative) $y=1$	twCRPS ($\delta=15$, $h=2$, illustrative) $y=12$	twCRPS ($\delta=15$, $h=2$, illustrative) $y=25$
$N(0, 5^2)$	1.25	9.21	22.18	0.003	0.378	9.98
$N(0, 10^2)$	2.38	7.48	19.40	0.054	0.324	9.34
t_3, scale 3	0.95	9.71	22.57	0.003	0.381	9.91

- Calm hour ($y = 1$): the sharp forecasts win; depeg hour ($y = 25$): the wide one wins. twCRPS downweights calm hours (0.003 vs CRPS 1.25) and concentrates the score on the depeg region
- Source: own computation for three illustrative forecasts; no market data

BA_L4_crps

Results: Distributional Accuracy

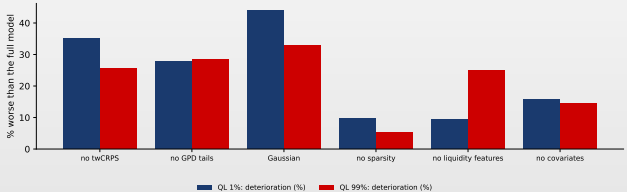


- Lowest loss on every metric; against the best benchmark: CRPS −6.6% (TimeXer), twCRPS −14.2% (TiDE)
- Against GARCH-t: CRPS 0.82 vs 2.21 (−63%), twCRPS 0.067 vs 0.106 (−37%)
- One-sided Diebold–Mariano tests with block-bootstrap intervals: all p-values below 0.05 against all five benchmarks
- Source: Chaffard et al. (2026), working paper, Tables 2–3

BA_L4_paper_t2

Calibration and Ablation

- Probability integral transform (PIT) $u_t = \hat{F}_t(y_t)$: mean 0.5022 (target 0.5), variance 0.0878 (target $1/12 = 0.0833$); Kolmogorov–Smirnov (KS) $D = 0.0288$, 48-hour block bootstrap p $= 0.079$
- Value-at-risk (VaR) exceedances: too many at the 2–5% levels in both tails, close to nominal at 0.5–1%; forecast ES more extreme than realised: conservative



- A Gaussian head loses 44% at the 1% quantile; dropping the liquidity features costs 25% at the 99% quantile, more than dropping all covariates
- Source: Chaffard et al. (2026), working paper; ablation from its Table 4

BA_L4_paper_ablation